Echo State Networks for Renewable Energy Forecasting

Samuel G. Dotson, Kathryn D. Huff

Dept. of Nuclear, Plasma and Radiological Engineering, University of Illinois at Urbana-Champaign
sgd2@illinois.edu

INTRODUCTION

Load following has given natural gas an economic edge over nuclear power because unlike nuclear, natural gas plants can follow grid demand and even shut off when renewable penetration makes the price of electricity go negative [1]. Advanced reactor designs, like Molten Salt Reactors (MSR), promise strong load following capabilities due to harder neutron spectra and faster Xe-135 burnup [2]. Unfortunately, the most mature MSR designs are at least a decade away from obtaining a commercial license in the United States. The climate crisis is too urgent to wait this long for nuclear power to become fully competitive with natural gas. Renewable energy has also challenged the base load electricity production that nuclear provides by introducing unpredictable variability in grid demand. Nuclear energy can be more economically feasible by relaxing the strong load following requirements with improved predictions of renewable energy production several hours or days in advance. If reactor operators knew in advance how much electricity will be produced by renewable energy they can slowly and accurately ramp reactor power to meet demand rather than operate continuously at full power and risk paying to export electricity. Thus improving nuclear energy's comptetitiveness against natural gas and strengthening nuclear's ability to couple with renewable energy. In this work we introduce Echo State Networks (ESN) as a preferred method for time series forecasting of chaotic and stochastic systems like electricity production from renewable sources.

BACKGROUND

The University of Illinois at Urbana-Champaign is an ideal model system for this work because of its diverse energy mix. Previous work has been done to characterize this energy grid and optimize the size of a nuclear reactor [3]. Due to the degree of wind penetration, the University is sometimes forced to sell electricity back to MISO at a loss because of overproduction from wind energy. Thus, a reliable prediction of electricity production from wind and other variable sources will reduce the likelihood of these events, specifically. In general, this prediction will also allow nuclear power plants to bid on the wholesale market, rather than act as price takers. Echo State Networks (ESN), a flavor of reservoir computing, are a modern machine learning algorithm that enables accurate short to medium term predictions. Pathak et. al used an ESN to predict the evolution of a chaotic system, a laminar flame front, up to seven Lyapunov times in the future [4, 5]. A Lyapunov time simply measures the timescale at which chaos makes initial predictions useless. The effect of chaos typically overwhelms conventional predictions after a single Lyapunov time, by definition. The Lyapunov time for a weather system is on the order of a few days but depends on regional environment. ESNs have also been used to forecast multivariate time series [6]. Echo state networks are unique among neural networks in their ease of implementation. However, the simplicity of their implementation is balanced by the need for carefully chosen hyperparameters for the desired task [7]. Electricity production from solar and wind are tightly coupled to regional weather. Electricity demand exhibits some seasonal regularity but is still subject to stochasticity. Combining accurate demand predictions with reliable renewable energy predictions will enable an artificially intelligent reactor operator to adjust power in a relaxed manner. Additionally, reservoir computing is relatively computationally inexpensive and fast to train. This is owed to its sparse network architecture [4, 5, 8].

METHODOLOGY

This work consists of three parts, data generation, hyperparameter search, model training, and prediction. To generate data needed for training our model we used the `TrainARMA` functionality in the `RAVEN` tool from INL to generate synthetic histories. These scenarios are passed to the ESN model for training and testing. The ESN is implemented with the open-source code `pyESN` [9].

Scenario Generation

In this work we use the scenario generation method described in Baker et. al using the `RAVEN` framework [10]. The algorithm for producing synthetic histories is:

1. Create a "typical year" from historical data.

2. Fit an Autoregressive Moving Average (ARMA) model using `TrainARMA`.

3. Use Monte Carlo sampling to generate synthetic histories from ARMA model.

We perform this process for three datasets: Electricity demand, wind generation, and solar generation. This allows us to train an ESN model to predict generation from wind and solar farms and to predict the net demand profile given by

$$D_{net} = (D_{total} - P_{wind} - P_{pv})_h \ \forall \ h \text{ in } [0, 8759] \qquad (1)$$

Where D_{total} is the total demand at a given hour, h, of a synthetic demand profile from the ARMA model.

Reservoir Computing Model

Unlike a typical feed forward neural network, an "echo state network" (also called a "liquid state machine") is a type of recurrent neural network that uses a single layer of many neurons called a "reservoir". The reservoir has an adjacency matrix A that is

1. sparsely populated

2. connected by uniformly random weights centered at zero

3. has a large number of neurons

A reservoir computer also satisfies the *echo state property* [4, 11]. This property ensures that a system's state has a decaying influence on future states (like an echo of sound or ripples on water). This property is satisfied in most cases when the spectral radius (the absolute value of the greatest eigenvalue of A)[11] is,

$$\rho(A) < 1. \qquad (2)$$

However, the echo state property can still be satisfied for a spectral radius greater than unity [7].
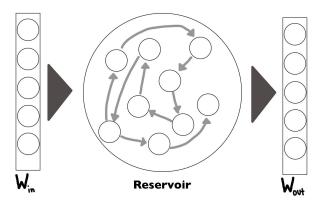
Fig. 1: A basic reservoir computer or echo state network.The connections in the reservoir are given by A

Figure 1 gives an visual representation of a basic ESN. An input vector of length K is mapped to the reservoir layer by an input weight matrix W_{in}. The state of the reservoir is mapped to an output layer of length N with an output weight matrix W_{out}. In this work, the input vector is a function of time, $u(t)$, and the output vector is the next state of the system, $u_p(t + \Delta t)$. Ideally, the difference between the prediction, u_p, and the actual, u_a, is minimized. During training, the output weight matrix is trained through backpropagation using a loss function like cross entropy [4, 12].

Hyperparameter Search

Due to the architecture of ESNs the weights and connections inside the reservoir do not need to be trained and, in our choice of implementation, cannot be. This dramatically reduces the training time because only the linear output layer needs to be trained. One drawback of this approach is its sensitivity to hyperparameters, which must be carefully chosen before running the network. Here, we perform grid searches over a variety of networks to establish which combination of hyperparameters minimizes the mean squared error of the model,

$$MSE = \frac{1}{N} \sum_{i}^{N} (\hat{y} - y_i)^2 \qquad (3)$$

where

$\hat{y}$ = the average value of the ouput.

RESULTS

Current results include the generation of time series data to be used in training the RC model. Shown in Figure 2, Figure 3, and Figure 4.
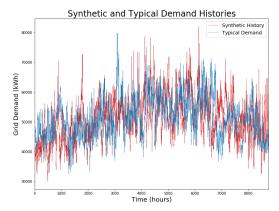

Fig. 2: The typical year of hourly grid demand in kW at UIUC.
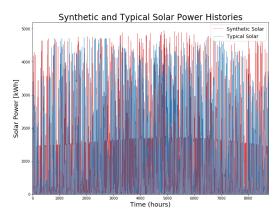

Fig. 3: The typical year and a synthetic year of hourly solar electricity generation in kWh per hour at UIUC. Data from [13]

The hyperparameters of the ESN used in Figure 5 and Figure 6 have not been optimized. In spite of this, preliminary predictions track reasonably well with grid demand. We conducted a single grid search for the an optimal combination of spectral radius (ρ) and noise injection (for regularization of reservoir neurons), Figure 7, following the recommendations from [7].

CONCLUSIONS

In this work we demonstrated the RAVEN tool's scenario generation capability. We have also demonstrated that Echo State Networks can powerfully predict the evolution of dynamic, chaotic, systems as others have [4, 5, 6]. Accurate predictions of chaotic systems, like wind energy production,
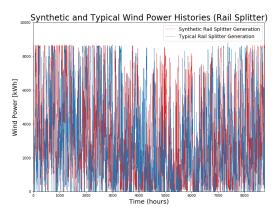
Fig. 4: The typical year and a synthetic year of hourly power produced by the UIUC wind power purchase agreement with Railsplitter Wind Farm.
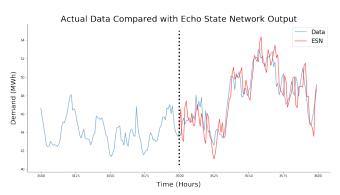

Fig. 5: A simple ESN with a prediction of 100 hours into the future. After training on 1000 hours of historical data.
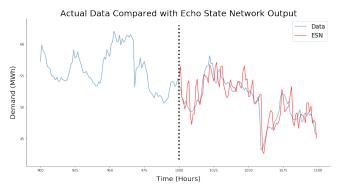

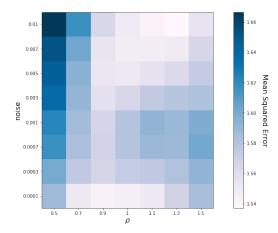Fig. 6: A simple ESN with a prediction of 100 hours into the future after training on 3500 hours of historical data.

Fig. 7: A grid search over a range of spectral radii and noise levels. The optimal set minimizes the mean squared error.

will enable nuclear power plants to maintain their economic feasibility. Knowledge, in advance, of renewable electricity generation will enable operators to ramp the reactor power within the constraints given by reactor physics. Future work will include identifying ideal input vectors, grid searches and random searches to tune hyperparameters, and averaging over several reservoirs, to achieve further improvements. ESNs will also be used to predict solar and wind energy production.

ACKNOWLEDGMENTS

REFERENCES

1. J. H. KEPPLER, C. MARCANTONINI, O. N. E. AGENCY, and O. FOR ECONOMIC CO-OPERATION {AND} DEVELOPMENT, *Carbon pricing, power markets and the competitiveness of nuclear power*, Nuclear development, Nuclear Energy Agency, Organisation for Economic Co-operation and Development.

2. A. RYKHLEVSKII, D. O'GRADY, T. KOZLOWSKI, and K. D. HUFF, "The Impact of Xenon-135 on Load Following Transatomic Power Molten Salt Reactor," in "Transactions of the American Nuclear Society," American Nuclear Society.

3. S. G. DOTSON, "Optimal Sizing of a Micro-reactor for Embedded Grid Systems," American Nuclear Society Annual Meeting 2020.

4. J. PATHAK, B. HUNT, M. GIRVAN, Z. LU, and E. OTT, "Model-Free Prediction of Large Spatiotemporally Chaotic Systems from Data: A Reservoir Computing Approach," **120**, *2*, 024102, publisher: American Physical Society.

5. A. WIKNER, J. PATHAK, B. HUNT, M. GIRVAN, T. AR-COMANO, I. SZUNYOGH, A. POMERANCE, and E. OTT, "Combining machine learning with knowledge-based modeling for scalable forecasting and subgrid-scale

closure of large, complex, spatiotemporal systems," **30**, *5*, 053111, publisher: American Institute of Physics.

6. F. M. BIANCHI, S. SCARDAPANE, S. LØKSE, and R. JENSSEN, "Reservoir computing approaches for representation and classification of multivariate time series," .

7. M. LUKOŠEVIČIUS, "A Practical Guide to Applying Echo State Networks," in G. MONTAVON, G. B. ORR, and K.-R. MÜLLER, editors, "Neural Networks: Tricks of the Trade: Second Edition," Springer, Lecture Notes in Computer Science, pp. 659–686.

8. S. VANNITSEM, "Predictability of large-scale atmospheric motions: Lyapunov exponents and error dynamics," **27**, *3*, 032101.

9. C. KORNDÖRFER, "pyESN," .

10. T. E. BAKER, A. S. EPINEY, C. RABITI, and E. SHITTU, "Optimal sizing of flexible nuclear hybrid energy system components considering wind volatility," **212**, 498–508.

11. M. LUKOŠEVIČIUS and H. JAEGER, "Reservoir computing approaches to recurrent neural network training," **3**, *3*, 127–149.

12. P. R. VLACHAS, J. PATHAK, B. R. HUNT, T. P. SAPSIS, M. GIRVAN, E. OTT, and P. KOUMOUTSAKOS, "Backpropagation algorithms and Reservoir Computing in Recurrent Neural Networks for the forecasting of complex spatiotemporal dynamics," **126**, 191–217.

13. ALSOENERGY, "University of Illinois Solar Farm Dashboard," Http://s35695.mini.alsoenergy.com/Dashboard/2a5669735065572f4a42454b772b714d3d.